

Improving Credit Card Fraud Detection using Generative Adversarial Networks (GAN)

6501 Capstone Group 4

Team Member:

Hao Ning, Jun Ying

Outline

- Problem Statement
- Exploratory Data Analysis (EDA)
- GAN Theory
- Base Model
- GAN & Performance Evaluation
- WGAN & Performance Evaluation
- WGAN-GP & Performance Evaluation
- GAN with Pre-Train Performance
- BEGAN & Performance Evaluation
- BAGAN & Performance Evaluation
- Conclusions

Problem Statement

- Credit Card Fraud Detection
 - Imbalanced dataset
 - European transactions within two days in September 2013
 - Kaggle Dataset: <https://www.kaggle.com/mlg-ulb/creditcardfraud>
- Oversampling & Undersampling
 - Biased towards the mainstream category
- Oversampling with GANs
 - How to evaluate the quality of generated data
- GANs Comparison
 - WGAN, WGAN_GP, BAGAN, BEGAN

EDA

- About Dataset: 284,807 transactions

-

```
Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',  
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',  
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',  
      'Class'],  
      dtype=object)  
  
   Time      V1      V2      V3  ...      V27      V28  Amount  Class  
0  0.0 -1.359807 -0.072781 2.536347 ...  0.133558 -0.021053  149.62     0  
1  0.0  1.191857  0.266151 0.166480 ... -0.008983  0.014724    2.69     0  
2  1.0 -1.358354 -1.340163 1.773209 ... -0.055353 -0.059752  378.66     0  
3  1.0 -0.966272 -0.185226 1.792993 ...  0.062723  0.061458  123.50     0  
4  2.0 -1.158233  0.877737 1.548718 ...  0.219422  0.215153   69.99     0
```

There are 30 features and 1 class (normal:0, fraud:1) & no null values.

Features V1 - V28 are the principal components obtained with PCA





- Imbalance Dataset: Only 0.173% data are fraud transactions.

-

```
The amounts of normal transactions (class 0) & fraud transactions (class 1)  
0    284315  
1      492
```

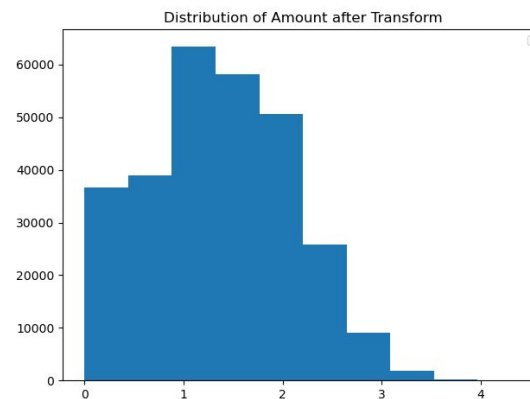

EDA

- Normal Transaction versus Fraud Transaction



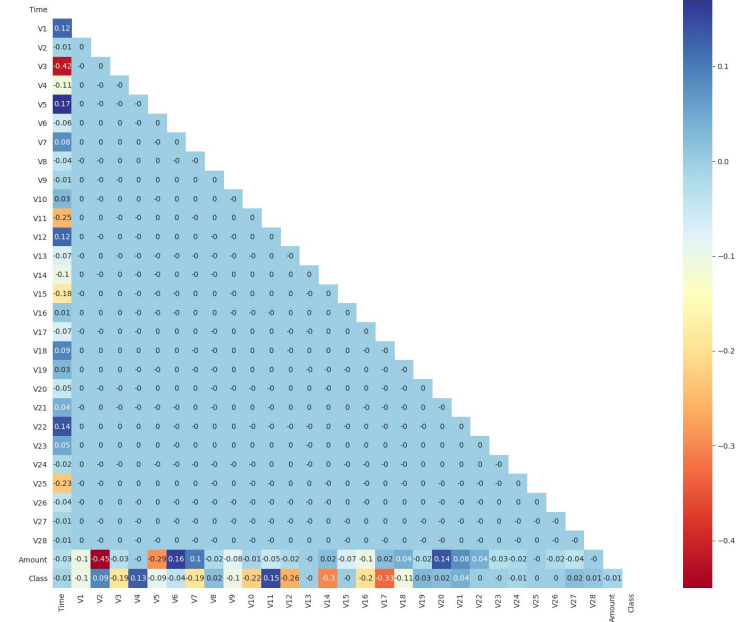

EDA

- Data Preprocessing
 - 'Time' is sorted in time series in seconds change to hours
 - $df['Time'] = (df['Time'].values / 3600)$
 - 'Amount' is larger than other features
 - $df['Amount'] = np.log10(df['Amount'].values + 1)$



EDA

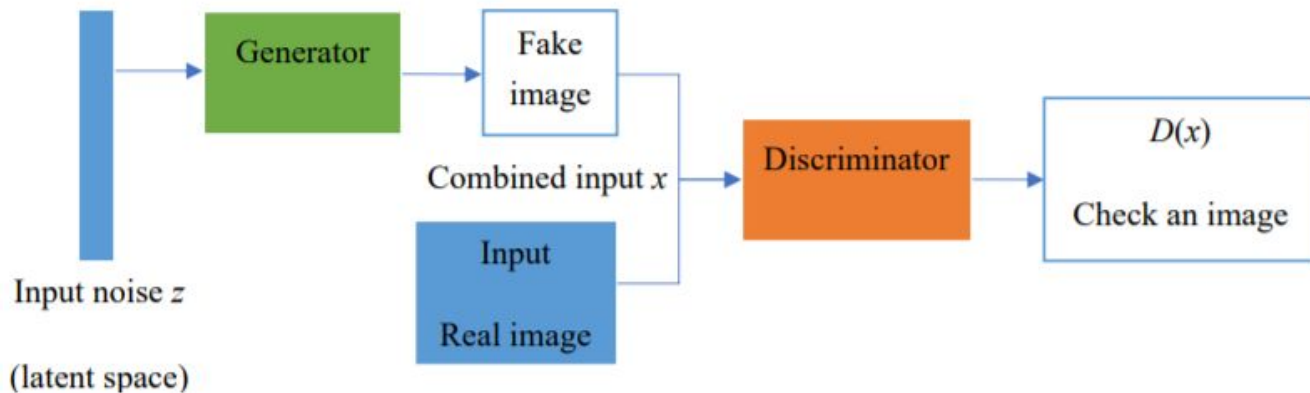
- Data Distribution & Correlation of the Features



GAN Theory

- Generative Adversarial Networks (GAN) is a cutting-edge technique of deep neural networks, which was first come up by Ian Goodfellow in 2014.


$$\min_D \max_G V(G, D) = \mathbb{E}_x \log(D(x)) + \mathbb{E}_z \log(D(G(z)))$$



GAN Theory

- GAN is a generative model in which two neural networks (generator and discriminator) compete and improve each other.

- Training Discriminator



- Training Generator




- Challenge

- How to verify that the generated data meets the requirements?

Base Model

- Train/Test Split & Stratified
 - Train 80% 227845 (394 fraud)
 - Test 20% 56962 (98 fraud)
- Random Under Sampling (RUS) & Random Over Sampling (ROS)
- GridsearchCV XGBoost
- Predict with best_params
- Test performance evaluation

	RUS 1 394 0 394	ROS/Base Model 1 227451 0 227451
Accuracy	0.954250	0.999491
Precision	0.033829	0.855670
Recall	0.928571	0.846939
F1 score	0.065280	0.851282
ROC AUC score	0.941433	0.923346



GAN

- Differences from vanilla GAN
 - Original GAN deal with images, but we are dealing with tabular data, tahn is removed
 - Batch normalization removed
- Noise = 32, len(features) = 30


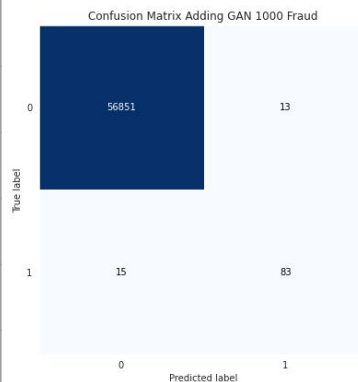
Generator Input	Generator output
(32, 1) Noise	(64,1)
(64, 1)	(128, 1)
(128, 1)	(256, 1)
(256, 1)	(30, 1) Features

Discriminator Input	Discriminator Output
(30, 1) Features	(256,1)
(256, 1)	(128, 1)
(128, 1)	(64, 1)
(64, 1)	(1, 1) label

GAN Performance

- Slight improvement, adding 1000 generated fraud
- Low spectrum of data
 - Mode collapse

	Base	GAN
Accuracy	0.999491	0.999491
Precision	0.855670	0.864583
Recall	0.846939	0.846939
F1 score	0.851282	0.855670
ROC AUC score	0.923346	0.923355



WGAN


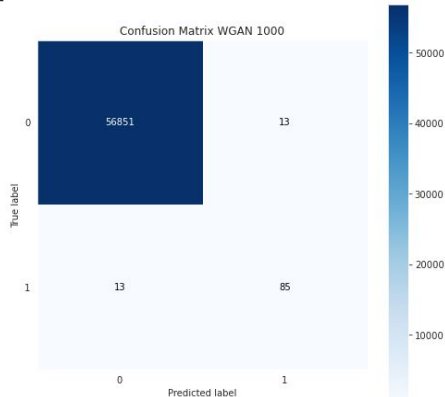
- Why WGAN
 - Prevent gradient vanishing & mode collapse in vanilla GAN
 - Evaluate the difference between real and generated samples with wasserstein distance, using a score rather than label
- Implementation compared to GAN
 - Loss function, no log
 - D: no sigmoid, and train D more than G
 - Clip the weight of D at (-0.01,0.01)
 - Use RMSProp

WGAN	D loss: $\tilde{V} = \frac{1}{m} \sum_{i=1}^m (D(x^i) - D(G(z^i)))$	G loss: $\tilde{V} = \frac{1}{m} \sum_{i=1}^m D(G(z^i))$
------	---	--

WGAN Performance

- Obvious improvement, adding 1000 generated fraud
- Wider spectrum of data
 - Overlap well with the original data
 - No mode collapse like GAN


	GAN	WGAN
Accuracy	0.999508	0.999544
Precision	0.864583	0.867347
Recall	0.846939	0.867347
F1 score	0.855670	0.867347
ROC AUC score	0.923355	0.933559



WGAN_GP

- Why WGAN_GP $|f(x_1) - f(x_2)| \leq K|x_1 - x_2|.$
 - A differentiable function is 1-Lipschitz function if and only if it has gradients with norm at most 1 everywhere
 - Hard to get gradients with norm at most 1 everywhere, interpolate between real and generated samples instead
 - Gradient penalty instead of weight clipping in WGAN


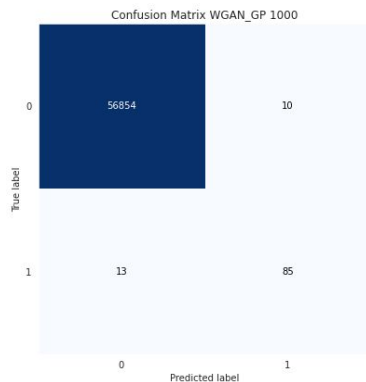
Gradient penalty: $\lambda(\|\nabla_{x_p} D(x_p)\| - 1)^2$



WGAN_GP Performance

- ADAM(lr=0.00001), batch_size=128, gp_lambda=5
- Obvious improvement, adding 1000 generated fraud
- Wider range & excellent overlap



	WGAN	WGAN_GP
Accuracy	0.999544	0.999596
Precision	0.867347	0.894737
Recall	0.867347	0.867347
F1 score	0.867347	0.880829
ROC AUC score	0.933559	0.933586



GAN with Pre-Train Performance


- Using the parameter of trained autoencoder
- Better than Original GAN, slightly worse than WGAN
- Wider spectrum of data
 - Overlap well with the original data

	GAN	GAN + AE
Accuracy	0.999508	0.999544
Precision	0.864583	0.882979
Recall	0.846939	0.846939
F1 score	0.855670	0.864583
ROC AUC score	0.923355	0.923373



BAGAN


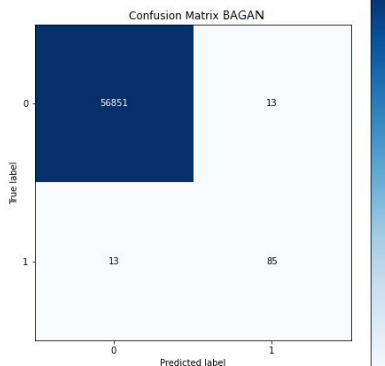
- Why BAGAN:
 - An augmentation tool
 - Learn from all data
- Compared to original GAN:
 - Discriminator output class label c or the label fake.
 - Coupling GAN and autoencoding techniques.



BAGAN Performance

- Performance better than GAN
- Generated fraud data same as the total dataset.
 - Learn too much from normal data
 - No similarity between these two classes

	GAN	BAGAN
Accuracy	0.999508	0.999544
Precision	0.864583	0.867347
Recall	0.846939	0.867347
F1 score	0.855670	0.867347
ROC AUC score	0.923355	0.933559



BEGAN


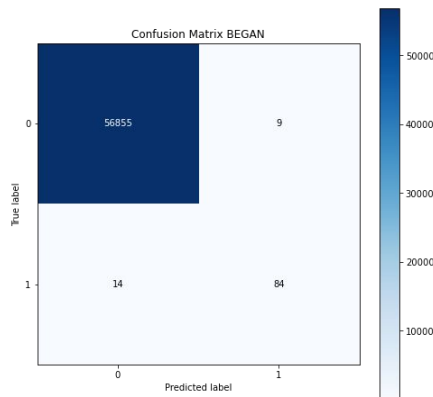
- Why BEGAN:
 - A GAN with simple yet robust architecture, standard training procedure with fast and stable convergence.
 - An equilibrium concept that balances the power of the discriminator against the generator.
 - Use Proportional Control Theory to maintain the equilibrium
 - Using a variable $k_t \in [0, 1]$ to control how much emphasis is put on $L(G(z_D))$ during gradient descent.

$$\begin{cases} \mathcal{L}_D = \mathcal{L}(x) - k_t \cdot \mathcal{L}(G(z_D)) & \text{for } \theta_D \\ \mathcal{L}_G = \mathcal{L}(G(z_G)) & \text{for } \theta_G \\ k_{t+1} = k_t + \lambda_k (\gamma \mathcal{L}(x) - \mathcal{L}(G(z_G))) & \text{for each training step } t \end{cases}$$

BEGAN Performance

- Performance better than BAGAN & WGAN
- Cover most of fraud data spectrum

	BAGAN	BEGAN
Accuracy	0.999544	0.999596
Precision	0.867347	0.903226
Recall	0.867347	0.857143
F1 score	0.867347	0.879581
ROC AUC score	0.933559	0.928492



Comparison of all Models

	Base	GAN	WGAN	WGAN_GP Best 1	GAN + AE	BAGAN	BEGAN Best 2
Accuracy	0.999491	0.999491	0.999544	0.999596	0.999544	0.999544	0.999596
Precision	0.855670	0.864583	0.867347	0.894737	0.882979	0.867347	0.903226
Recall	0.846939	0.846939	0.867347	0.867347	0.846939	0.867347	0.857143
F1 score	0.851282	0.855670	0.867347	0.880829	0.864583	0.867347	0.879581
ROC AUC score	0.923346	0.923355	0.933559	0.933586	0.923373	0.933559	0.928492

Conclusions

- GANs work well with tabular data with proper modification
 - Use data distribution to check the quality of generated data
- Vanilla GAN slightly improved fraud detection
 - Low spectrum data
- Improved GAN models showed better fraud detection performance
 - Wider range & better overlap
- **WGAN_GP** and **BEGAN** perform best among all GANs
- Using GAN as an oversampling strategy has great potential in credit card fraud detection and extremely imbalanced dataset

References

1. GAN [Generative Adversarial Networks](#)
2. Wasserstein GAN [\[1701.07875\] Wasserstein GAN](#)
3. Improved Training of Wasserstein GANs <https://arxiv.org/pdf/1704.00028>
4. BAGAN [\[1803.09655\] BAGAN: Data Augmentation with Balancing GAN](#)
5. BEGAN: Boundary Equilibrium Generative Adversarial Networks
[\[1703.10717\] BEGAN: Boundary Equilibrium Generative Adversarial Networks](#)
6. <https://github.com/eriklindernoren/Keras-GAN>
7. <https://github.com/IBM/BAGAN>

Thank you!

Q&A

Supplemental Materials

- JS Convergence vs Wasserstein
 - Loss function of GAN is equivalent to JS convergence (D at optimality)
 - JS is always $\log 2$ if two distributions do not overlap (0 gradient)

